

Reading Reference: Textbook 1 Chapter 2

# INTRODUCTION TO UNIX PROCESS

Tanzir Ahmed  
CSCE 313 Spring 2020

# Process Definition

- Process is an instance of a running programming. This happens through its **Abstraction**
  - *Also defined as “A Program in Action”*
  - *Provided by the OS and used by the program*
- To the OS, Process is a **data structure (and more)** representing a running program
  - *Used to save a program’s state*
  - *A program can be kicked out and restored using it*
- To the program, **Process** is a view of the entire system – memory, CPU, disk and all I/O devices
  - *Whatever the program needs, Process has it*
  - *Example: File handles, network socket etc. are kept inside the process*
  - *This view of the system is also called **machine state***

# Process Data Structures

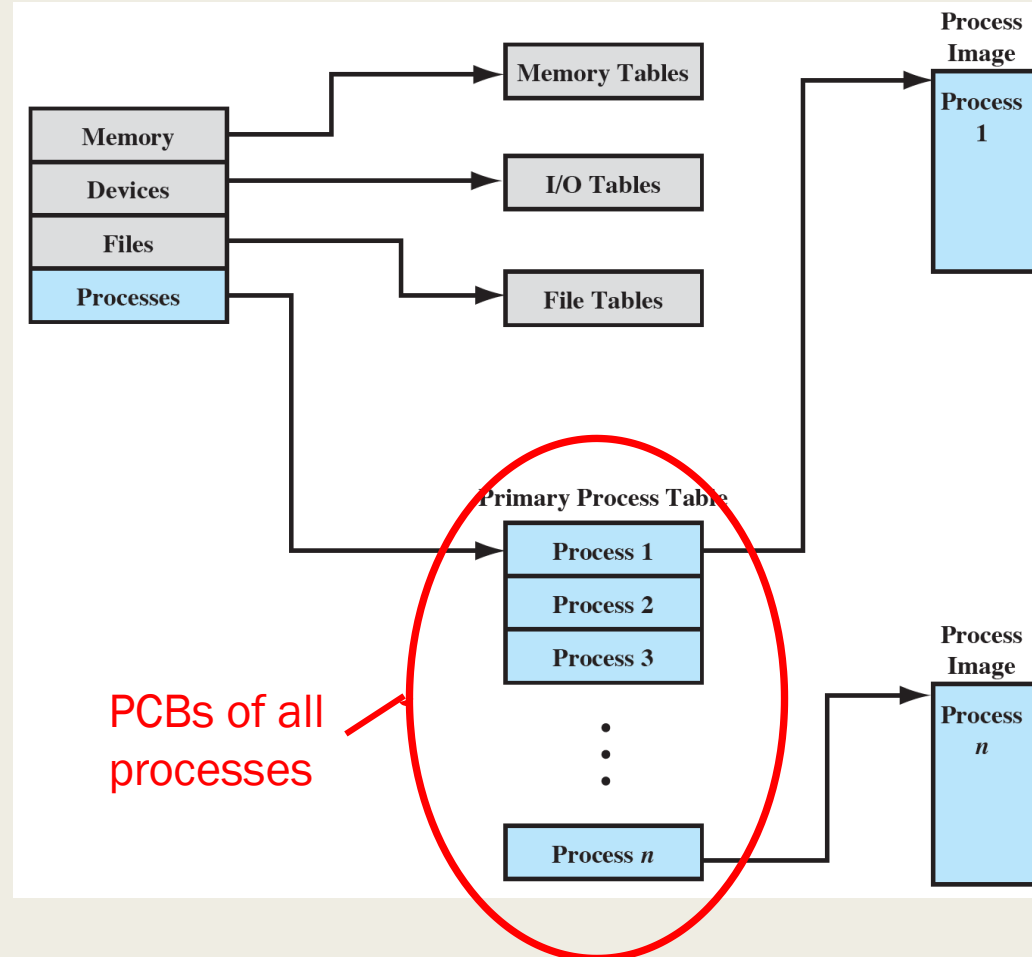
- How does the OS represent a process in the kernel?
  - Using a data structure called *Process Control Block (PCB)*
  - The PCB is where the OS keeps of a process' hardware execution state (PC, SP, regs, open file handles etc.) when the process is not running
    - This state is everything that is needed to restore the “kicked out” process in the same hardware configuration – as if nothing has happened
    - Required for transparent “kickout” and “bringing back in”
- The OS maintains a list of all Processes that are currently running (or some other relevant state)
  - Sometimes the OS needs to enumerate through processes
- Note that program data is not same as PCB
  - PCB is more like a process's metadata
  - Program's data (variables, allocated memory) and code are kept elsewhere (i.e., address space)

# Process Control Block (PCB)

process identification (use: to <b>locate</b> a processes)	<ul style="list-style-type: none"><li>• process id</li><li>• parent process id</li><li>• user id</li></ul>
processor state information (use: to <b>restore</b> a processes as it was)	<ul style="list-style-type: none"><li>• register set</li><li>• All general regs</li><li>• Specials (e.g., PC, SP, EFLAGS)</li><li>• condition codes</li><li>• Overflow, jump to take or not</li><li>• processor status</li></ul>
process control information (use: to <b>treat/run</b> a processes appropriately)	<ul style="list-style-type: none"><li>• process state</li><li>• scheduling information</li><li>• event (wait-for)</li><li>• memory-mgmt information</li><li>• owned resources (e.g., list of opened files)</li></ul>

# OS's Internal Tables

- An OS keeps a lot of information in the main memory
- In summary, much of this info is about:
  - *Resources (i.e., devices)*
  - *Running programs/processes*

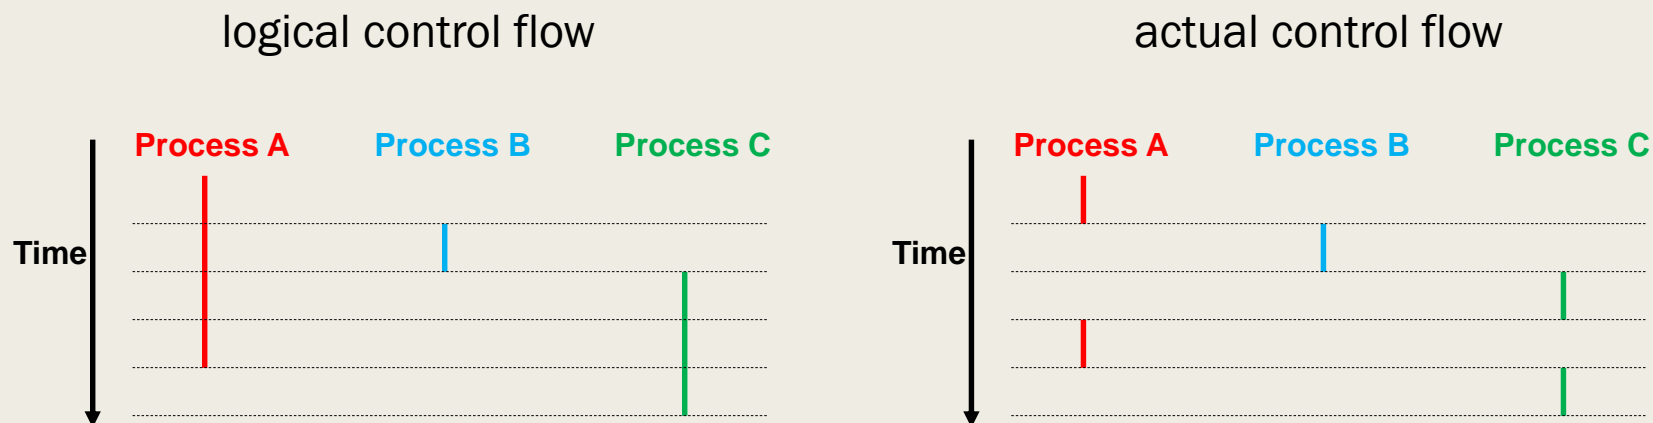


# Abstraction Mechanism

- Process provides each program with two key abstractions:
  - *Logical control flow for CPU virtualization*
    - Each program seems to have exclusive use of the CPU
  - *Private address space for Memory virtualization*
    - Each program seems to have exclusive use of main memory
- How are these illusions maintained?
  - Process executions *interleaved* (multitasking)
  - Private address spaces managed by virtual memory system (will describe that in a bit)

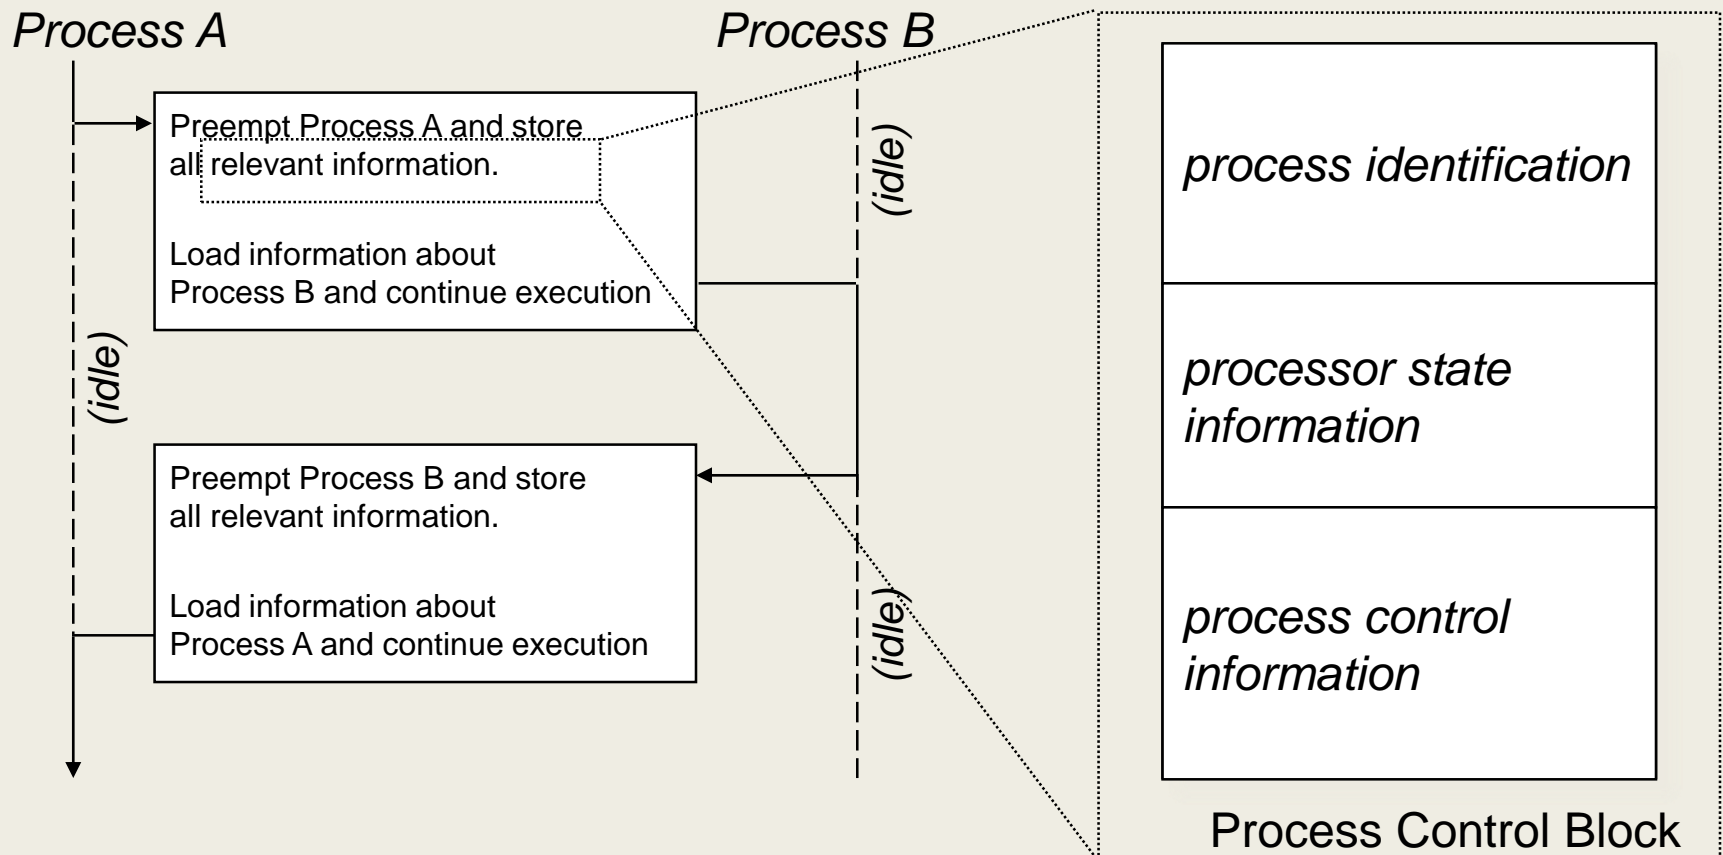
# Logical Control Flows

- Each process has its own logical control flow, which can be far from reality



- Control flows for concurrent processes are physically disjoint in time (except on multi-core machines)
- However, we can think of concurrent processes as running in 'parallel' with each other

# Logical Control Flow using Context Switch



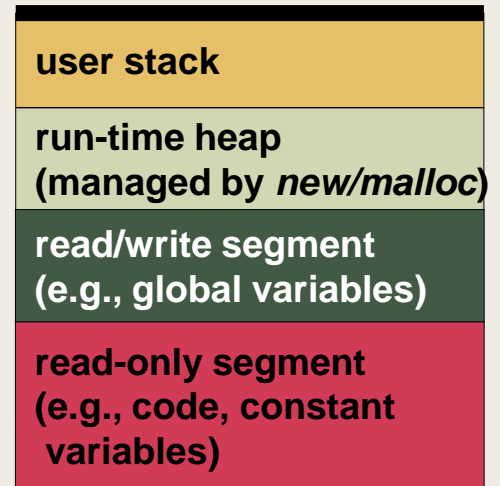
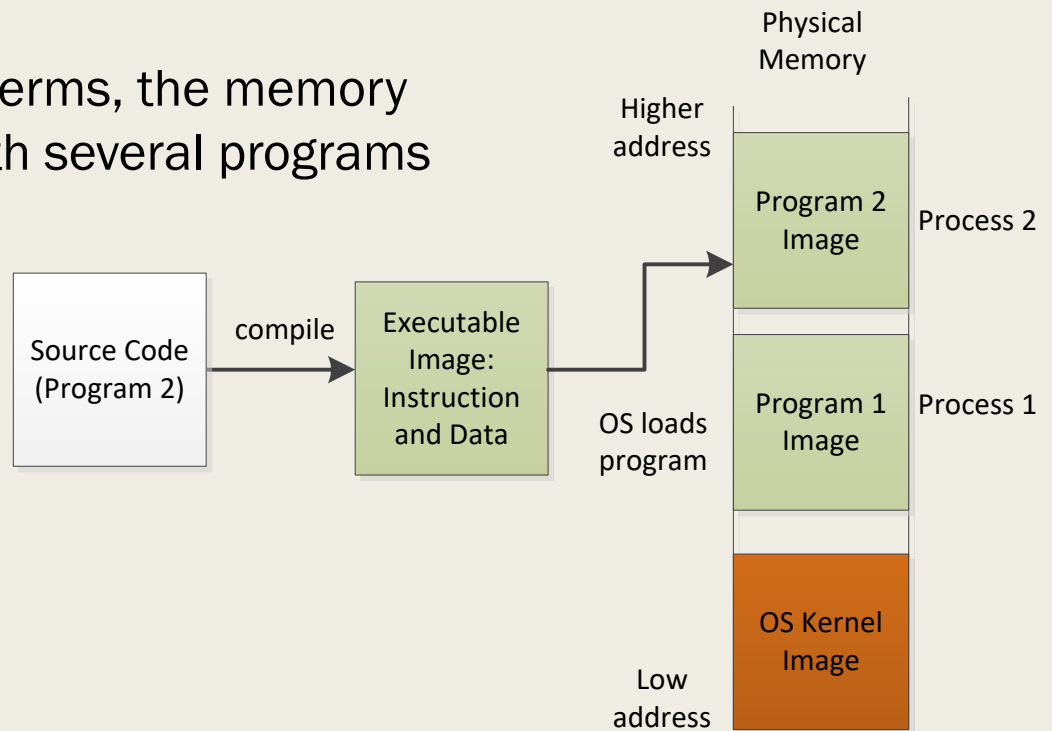


# Private Address Space

- First, every program needs to store and use some data
  - *Declaring variables, allocating memory etc.*
  - *A program's instructions also need to store*
- These data are kept in the disk, but need to be loaded to main memory before running the program
  - *Now, how and where do we load these data?*
  - *Answer: the “Address Space” of the corresponding Process*
  - *The address space is also “private” (i.e., isolated from other processes)*
- Let us see how it works

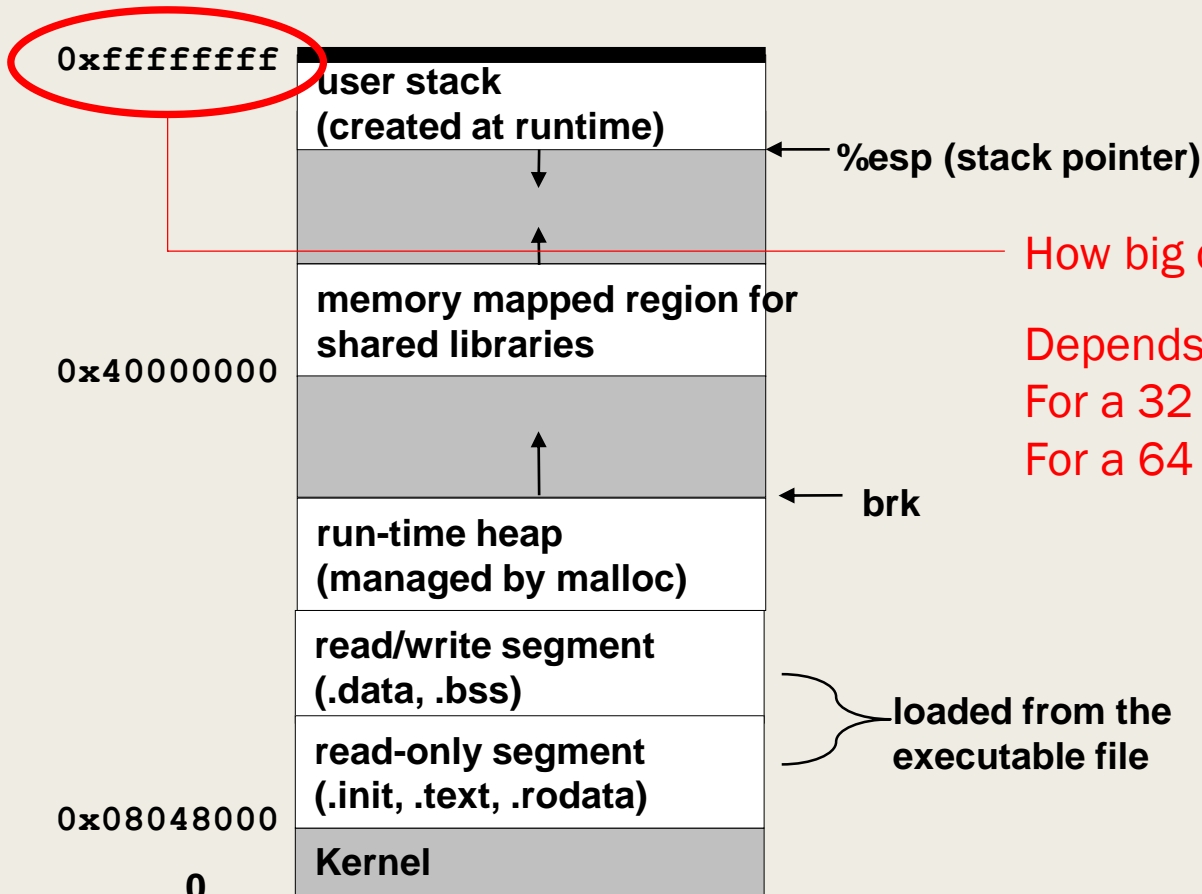
# Address Space

- A program's data can be divided into different parts; therefore they are loaded as groups into memory
- Therefore, in very simple terms, the memory looks like the following with several programs loaded:



# Private Address Space Illusion

- But each process is made to believe that the memory looks like the following – thanks to Virtual Memory:




# Question

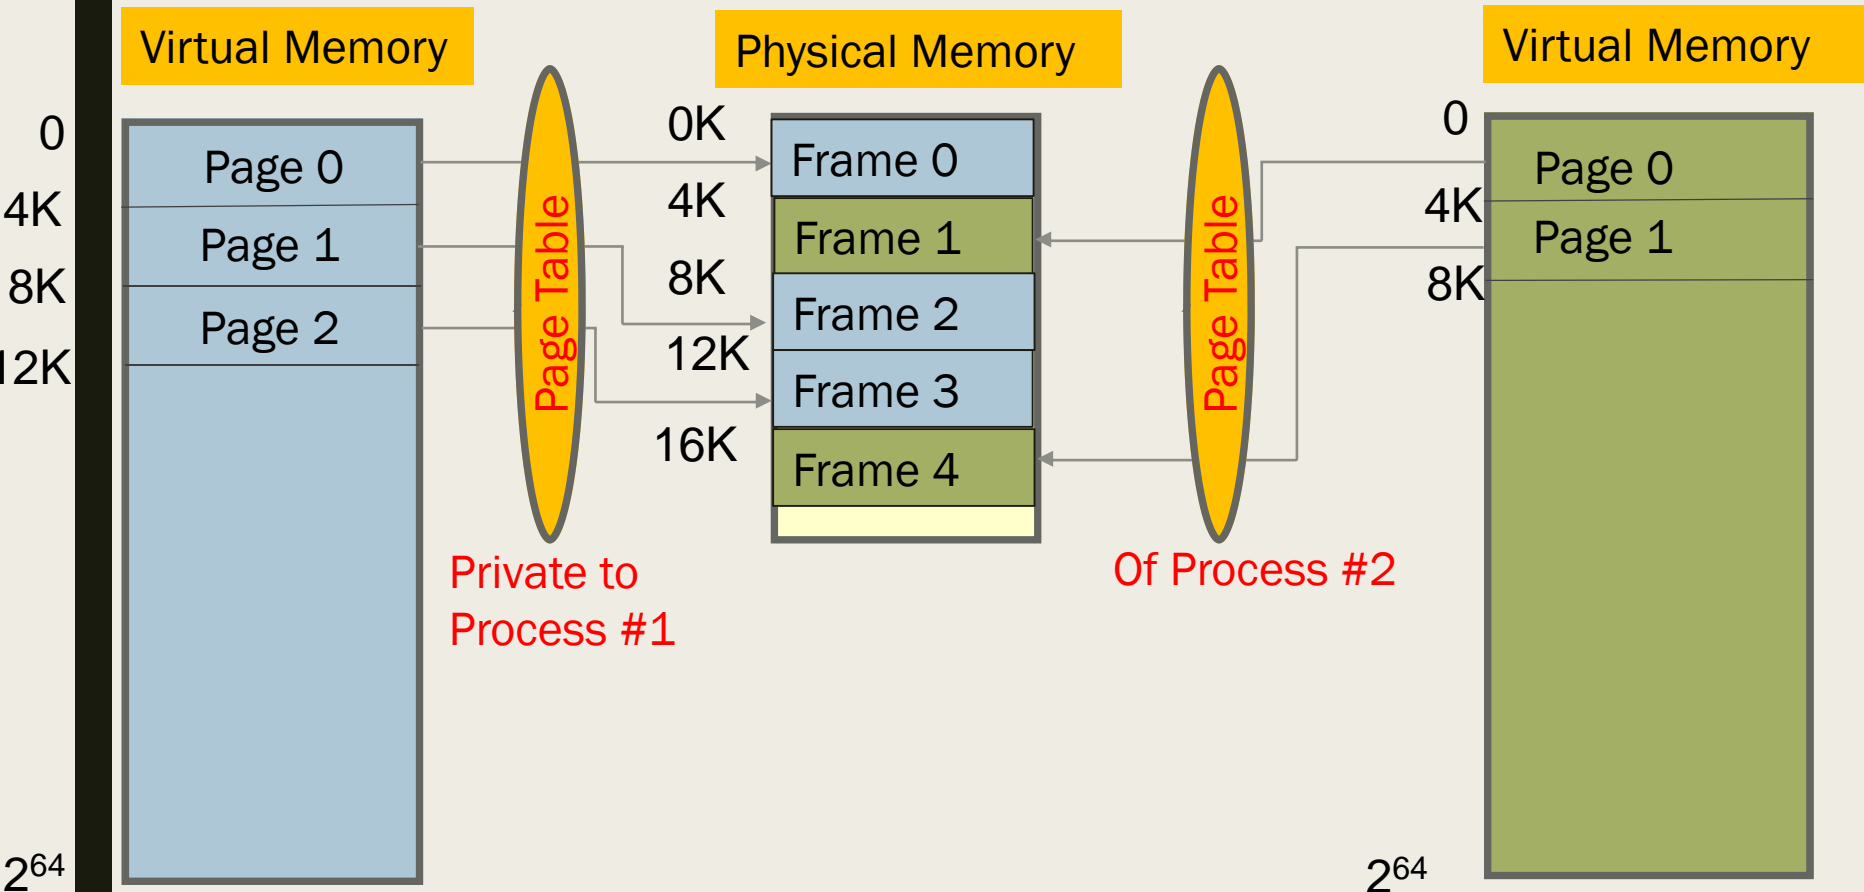
- Is a process image really  $2^{32} = 4\text{GB}$  long?
  - *It would be nice for us, programmers*
  - *But typically, physical memory is quite limited*
  - *Virtual memory in action*
- What is then Virtual Memory?
  - *For that, we now need to a little understanding on how memory is organized in a system*

# Memory Organization

- Memory is logically divided into *pages*, which are fixed in size and aligned regions of memory
  - *Typical size: 4KB/page*
- But pages are associated to a process only when necessary (i.e., read or written)
  - *When not necessary, they are evicted to the disk*
  - *Thus a running process can be stored back to disk again!!!!*
- The mechanism of such **lazy-allocation** is through **Page Fault**
  - *If a non-existent page is accessed (R/W), a page fault happens and fault-handler allocates the required page in physical memory*



# Mapping from Virtual to Physical Memory



# Summarizing Virtual Memory

- The private address space of a process is made up of pages
  - *These pages can be spread according to the wish of the kernel*
  - *Contiguous memory in processes' view are not necessarily contiguous in physical memory – they can be physically scattered, but **stitched** together by Virtual Memory system*
- Because memory is scarce, the whole private address space does not need to be allocated all the time:
  - *Actual pages can be allocated/mapped only when needed (i.e., read or written) through **page faults***
  - *Even allocated but inactive pages can be swapped back to disk to make room for more active pages*
- Memory accesses can be slowed down by Page Table accesses
  - *Each address must be translated to physical address by looking up in the process's page table, which is also in memory*
  - *Thus each memory access is actually **2 memory accesses**: 1 for page table, another for the actual memory access*
  - *There is a cache called Translation Lookaside Buffer (TLB) which stores popular translations – thus hiding the double latency*

# Virtual Memory – An Interesting Video

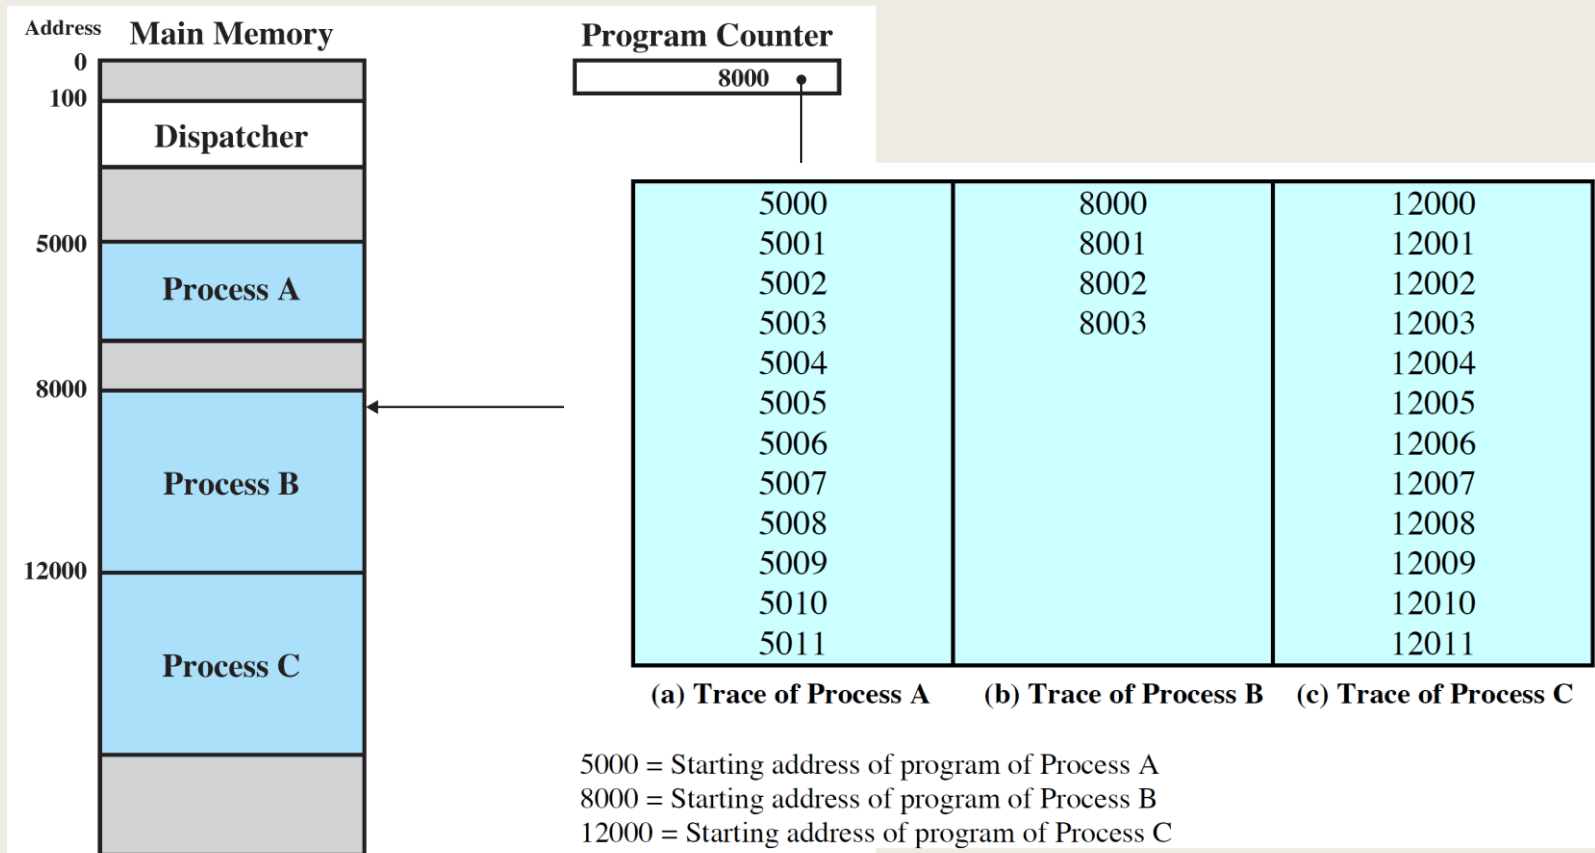
- Please watch this video, which is very informative, yet small:

<https://www.youtube.com/watch?v=qIH4-oHnBb8>



# Process States - An Example

- Assume the following processes A, B, C are loaded in memory
  - Assume we are not using Virtual Memory in this case
  - The processes are completely loaded in memory



# Process Trace

- **Trace:** The sequence of instructions that execute for a process

5000	8000	12000
5001	8001	12001
5002	8002	12002
5003	8003	12003
5004		12004
5005		12005
5006		12006
5007		12007
5008		12008
5009		12009
5010		12010
5011		12011

1 5000  
2 5001  
3 5002  
4 5003  
5 5004  
6 5005

-----Timeout

7 100  
8 101  
9 102  
10 103  
11 104  
12 105

13 8000  
14 8001  
15 8002  
16 8003

----- I/O Request

17 100  
18 101  
19 102  
20 103  
21 104  
22 105

23 12000  
24 12001  
25 12002  
26 12003

27 12004  
28 12005

-----Timeout

29 100  
30 101  
31 102  
32 103  
33 104  
34 105

35 5006  
36 5007  
37 5008  
38 5009  
39 5010  
40 5011

-----Timeout

41 100  
42 101  
43 102  
44 103  
45 104  
46 105

47 12006  
48 12007  
49 12008  
50 12009  
51 12010  
52 12011


-----Timeout

# Process Trace Discussion


- Dispatcher is show in blue shaded box
  - *The same dispatcher code runs between any 2 processes*
- The figure shows several (52) instruction cycles from the 3 processes
- The trace starts with process A by overwriting PC with 5000, which is the first instruction of A
- The OS allows exactly 6 instruction before the timer fires
  - *i.e., kicks out the currently running process*
  - *This prevents programs from monopolizing the CPU*
- CPU goes from A to B after A is kicked out for the timer
- B gets kicked out because of requesting I/O
- Then the CPU alternates between A and C because B is still waiting for the I/O

# Process States

- Let us start with elementary 2-state model
- This queue is some sort of a priority queue,
  - *priority is based on some scheduling metric*
- Typical content of the Queue:
  - *Pointer to PCB*
- **Limitation:**
  - *Does not handle I/O operation well*
  - *Can bring a process that is still waiting on I/O*
  - *The process will stay idle*



(a) State transition diagram

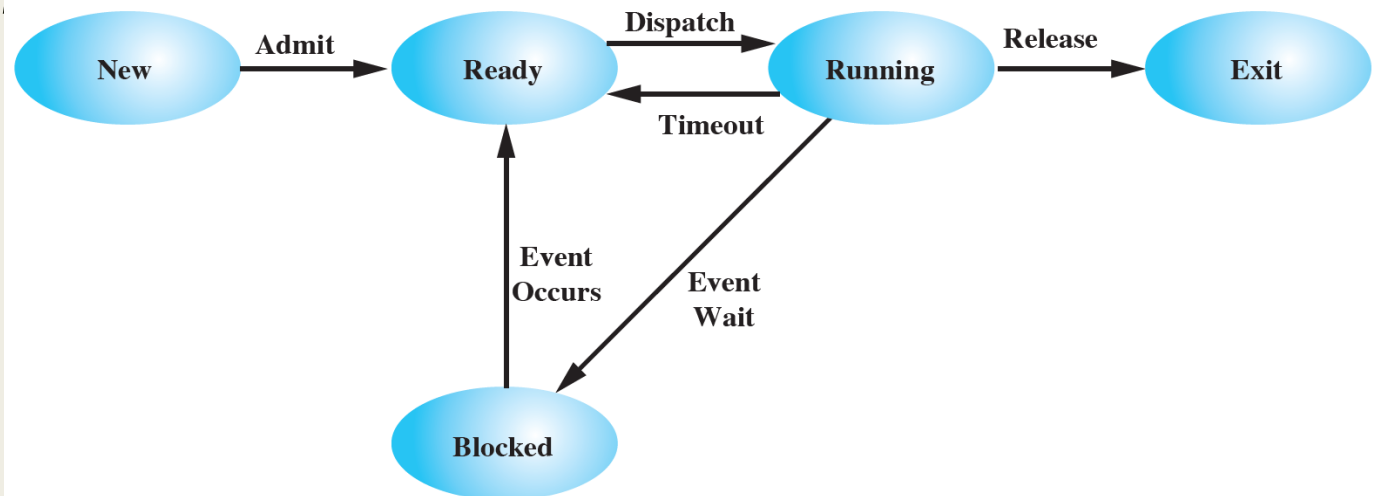


(b) Queuing diagram

Figure 3.5 Two-State Process Model

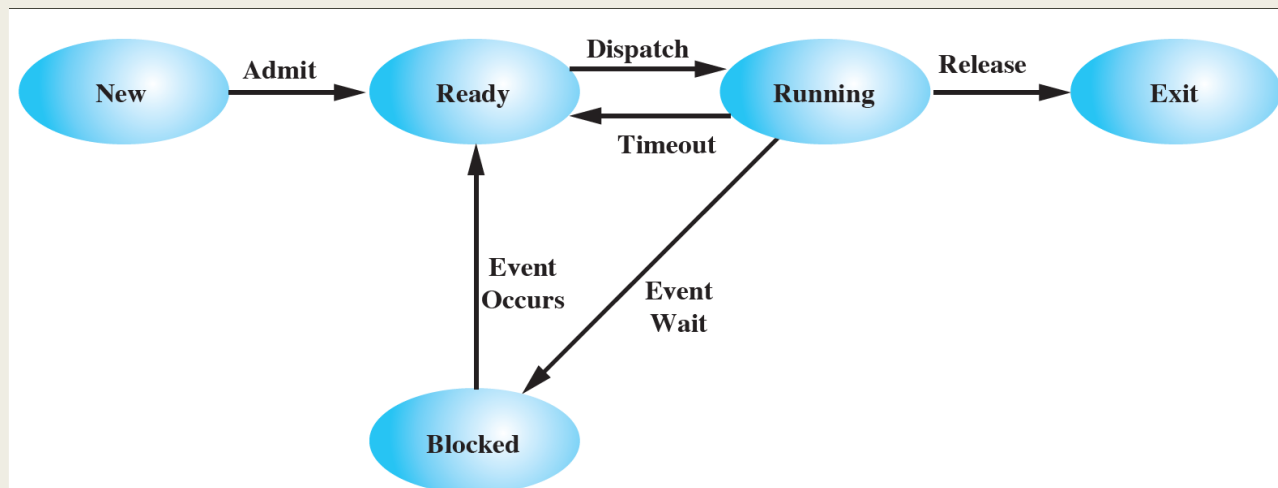
# Refined Process State Model

- It has 5 states
  - Newly added states: “New”, “Blocked”, “Exit”
- Addition of “Blocked” state is obvious
  - Avoids bringing the process back to ready queue before I/O operation or the event actually finished
- “New” state is when the PCB is created, but the process is not loaded in memory yet
  - Either because it is the usual delay
  - Or

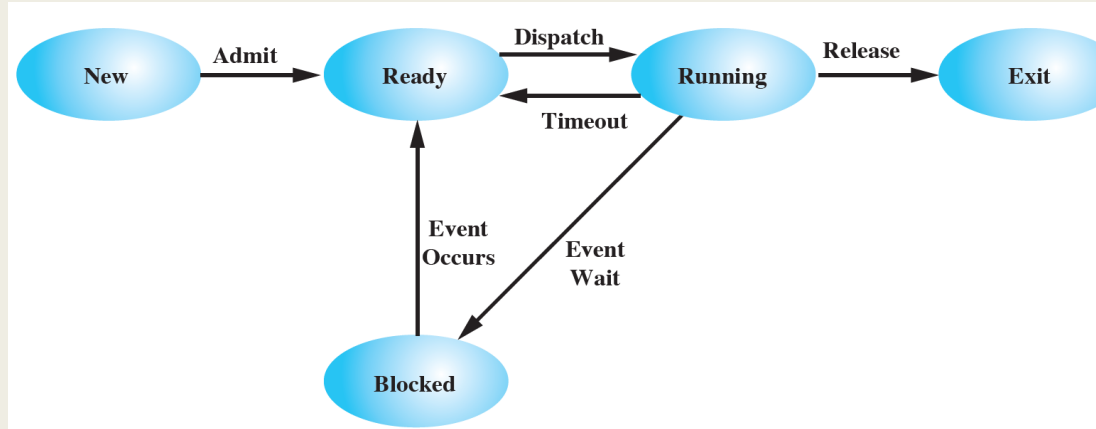


# Refined Process State Model (2)

- Exit State: The process is removed from the list of executable processes, but it is held by the OS for collecting some information about it
  - *E.g., an accounting program collecting information about all processes*
- New State: Is used for memory management
  - *There might a limit on how many processes can be there in the main memory*
  - *The process's PCB is in memory, but not its executable image*



# Now, the Transitions



- New->Ready:
  - *When there is room for a new process in memory*
- Ready->Running:
  - *When the scheduler picks this process*
- Running->Exit:
  - *Normal or due to some unavoidable/unrecoverable error (e.g., segmentation fault, divide by 0)*
- Running->Ready:
  - *Timer fired*
  - *A low-priority process is running in CPU and a higher priority process got unblocked for I/O finish*

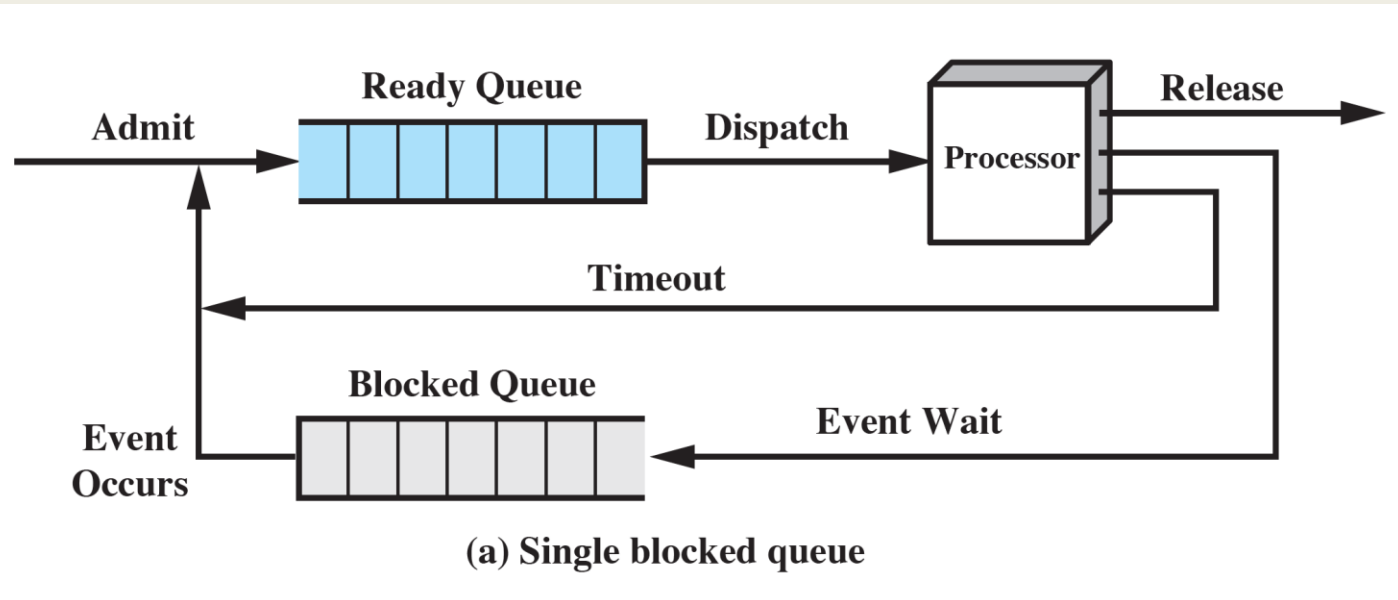
# Transitions

- Running->Blocked
  - *I/O or other event request that are not ready, or would take time*
- Blocked->Ready
  - *Event on which the process is waiting has finished*
- Ready->Exit
  - *Parent process terminates child before it could even run*
- Blocked->Exit
  - *Parent process killed the child while it was waiting*



# Queueing Model for Proc. States

- To implement the 5-state model, we will now need 2 queues



- However, 2-queue model is also not enough
  - *Does not make sense to make a file request from disk and URL request behind the Network card wait in the same queue*
- In reality, each I/O device, each lock, and thus each event needs a separate wait queue

# Queueing Model (2)

