

Reading Reference:

Textbook1: Chapter 4 (Section 4.1-4.4)

CONCURRENCY AND THREADS

Tanzir Ahmed
CSCE 313 Spring 2020

Today's Conversation

■ Threads

- *A bit complex topic but central to our understanding of modern computer systems (HW and SW)*
- *We will build concepts incrementally and tie the picture together at the end*

Adapted from contemporary courses in OS/Systems taught at Berkeley, UW, TAMU, UIUC, and Rice. Special acknowledgment to Profs Gu/Bettati/Tyagi at TAMU, Culler and Joseph at Berkeley

Why Processes & Threads?

Goals:

- **Multiprogramming:** Run multiple applications concurrently
- **Protection:** Don't want a bad application to crash system!



Solution:

- Virtual Machine abstraction: give process illusion it owns machine (i.e., CPU, Memory, and IO device multiplexing)
- Process: unit of execution and allocation



Challenge:

- Process creation & switching expensive
- Need concurrency within same app (e.g., web server)



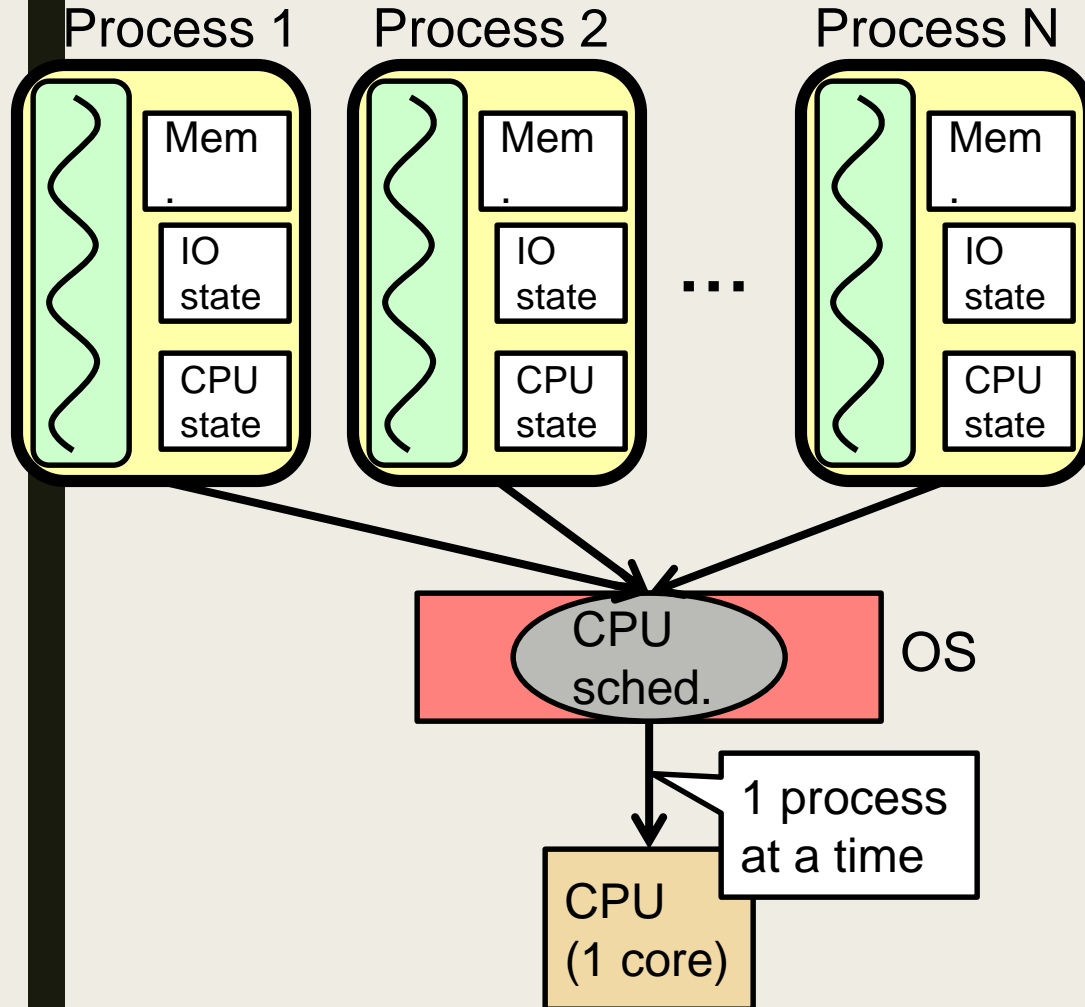
Solution:

- Thread: Decouple allocation and execution
- Run multiple threads within same process

Motivation for Threads

- Process Context Switch has huge overhead. Requires switching:
 - *CPU state (PC, PSW, all registers etc)* - fast
 - *Memory Address Space* - **SLOW**
- **Assumption: memory access without cache is slow**
- A Memory Address Space comes with:
 - *Virtual Memory System data*
 - Page tables (mapping) and Translation Look-aside Buffer (TLB) (i.e., working as cache for mapping)
 - *Actual Memory Content*
 - Dirty (modified) pages that are updated only in cache, not in memory
 - Need to be updated in the memory before switch
- All of these become obsolete after a switch
 - *All caches (L1, L2, ..), TLB become **COLD***
 - *Require time to **WARM UP** as the new process runs*

Processes Context Switch



- Process Switch overhead: **high**
 - CPU state: **low**
 - Memory/IO state: **high**
- Process creation: **high**
- Protection
 - CPU: **yes**
 - Memory/IO: **yes**
- Sharing overhead: **high** (involves at least a context switch)


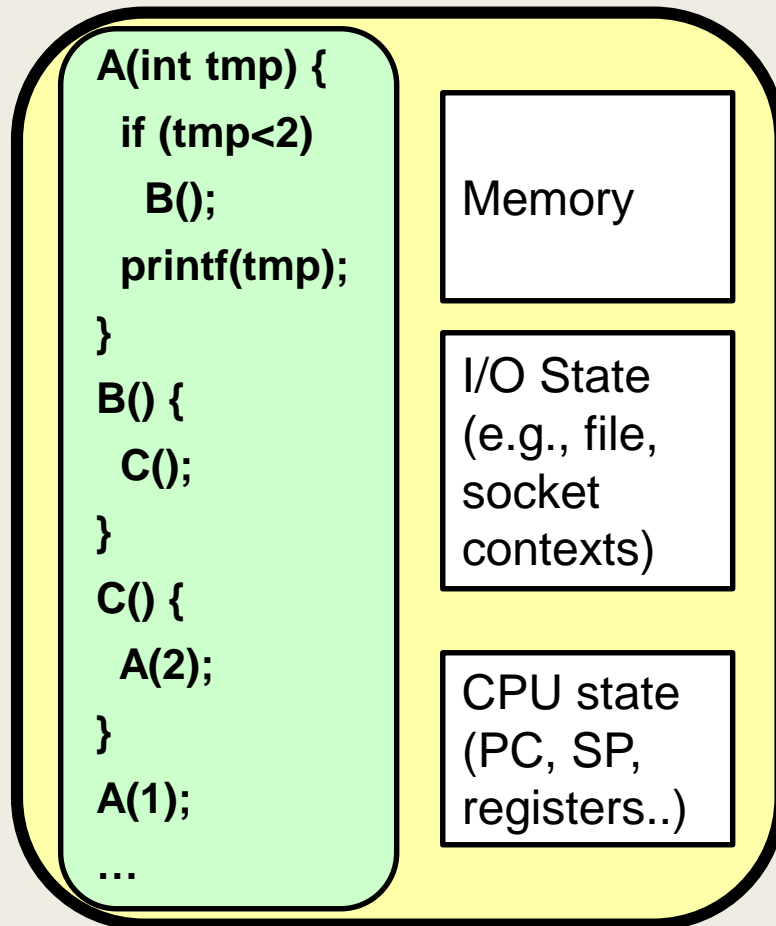
Thread Context Switch

- Still requires:
 - *The kernel scheduler to run, that overhead stays*
 - *Save the current CPU context, and load the next one*
- But, the old address space stays. Thus, no **WARM UP** required for:
 - *Caches (L1, L2, L3, ...)*
 - *TLB (cache for memory mapping)*

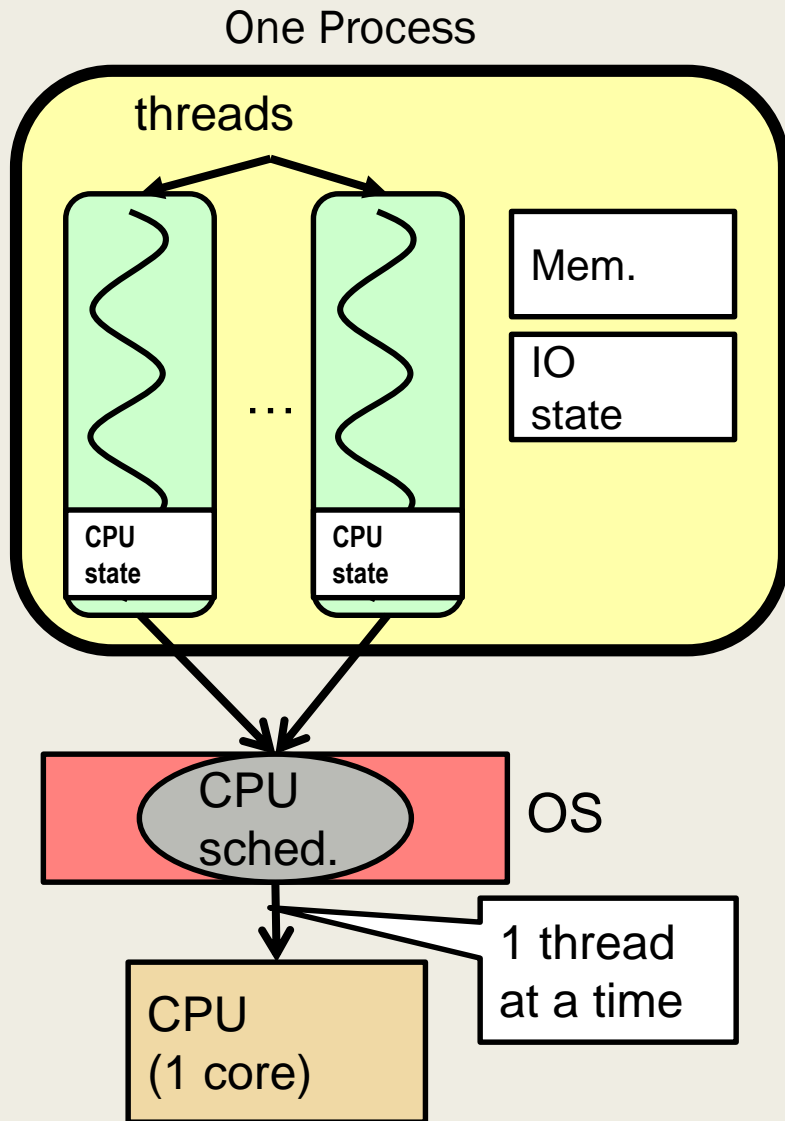
Saves time
- **Result: Thread Context Switch is much Faster than Process Context Switch**

Putting it together: Process

(Unix) Process



Threads Context Switch



- Thread Switch overhead: **low**
 - Only CPU state switched
- Thread creation: **low**
- Protection
 - CPU: **yes**
 - Memory/IO: **No**
- Sharing overhead: **low**
 - No context switches
 - Only low-overhead thread-switches

Threads in Use

- **Operating systems** need to be able to handle multiple things at once (MTAO)
 - *processes, interrupts, background system maintenance*
- **Servers** need to handle MTAO
 - *Multiple connections handled simultaneously in a Web Server*
- **Parallel programs** need to handle MTAO
 - *To achieve better performance*
- **Programs with user interfaces** often need to handle MTAO
 - *To achieve user responsiveness while doing computation*
- **Network and disk bound programs** need to handle MTAO
 - *To hide network/disk latency*

There's a Problem!!!

- Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A

x = 1;

Thread B

y = 2;

- However, What about (Initially, y = 12):

Thread A

x = 1;

x = y+1;

Thread B

y = 2;

y = y*2;

- *What are the possible values of x?*

Thread A

x = 1;

x = y+1;

Thread B

y = 2;

y = y*2

x=13

Problem is at the lowest level

- Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A

`x = 1;`

Thread B

`y = 2;`

- However, What about (Initially, $y = 12$):

Thread A

`x = 1;`

`x = y+1;`

Thread B


`y = 2;`

`y = y*2;`

- *What are the possible values of x ?*

Thread A

Thread B



`x=5`

Problem is at the lowest level

- Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A

`x = 1;`

Thread B

`y = 2;`

- However, What about (Initially, $y = 12$):

Thread A

`x = 1;`

`x = y+1;`

Thread B

`y = 2;`

`y = y*2;`

- *What are the possible values of x ?*

Thread A

Thread B



`x=3`

Race Condition is the name!!

- **Race condition:** Output of a concurrent program depends on the **order of operations** between threads
- Cannot make any assumptions about relative speed of threads (i.e. interleaving is a given)
- **Non-determinism** is omnipresent:
 - *Scheduler's decision depends on many factors*
 - *Processor architecture (e.g., variable clock rate, out-of-order execution)*

Race Condition for Instruction Set

- Simple threaded code (assume $x=0$)

Thread1

$x=x+1$;

Thread2

$x=x+1$;

Compiler Generated (because it can only use things from the Instruction Set):

load r, x

add r, r, 1

store x, r

Pre-emption

values of x can be 1 or 2 depending on the order of execution

load r1, x
add r1, r1, 1
store x, r1

load r2, x
add r2, r2, 1
store x, r2

$X=2$

load r1, x

add r1, r1, 1
store x, r1

load r2, x
add r2, r2, 1
store x, r2

$X=1$

Race Condition for Out-of-Order Execution

- Most architectures support (I should say “require”) this feature
 - *Pipelined processors cannot achieve peak performance without it*
- The idea is simple: compilers may reorder “unrelated” instructions like the following:

```
//global variables
bool done = false;
int data = -1;

Thread1 (){
    // takes a long time
    data = longfunction();
    done = true;
}
```



reorder

```
//global variables
bool done = false;
int data = -1;

Thread1 (){
    done = true;
    // takes a long time
    data = longfunction();
}
```

Out-of-Order Execution causing Race Condition

- The problem arises when there is an inter-thread dependency
- Because of thread 2's wait, the 2 lines in Thread1() are no longer independent
- However, the compiler has no way to tell that!!!

```
// thread 1's function,  
// REORDERD  
Thread1 (){  
    done = true;  
    data = longfunction();  
}
```

```
// thread 2's function  
Thread2 (){  
    while (!done); // wait  
    int newdata = compute (data);  
}
```


Thread API

- We will discuss thread API from C++11 which is cross-platform
 - *But for many other things, we will still use Linux for our PAs*
- Here is an example:

```
#include <iostream>
#include <thread>
#include <unistd.h>
using namespace std;

void foo() {
    sleep(3);
    cout<<"foo done"<<endl;
}

void bar(int x){
    sleep(1);
    cout <<"bar done"<<endl;
}
```

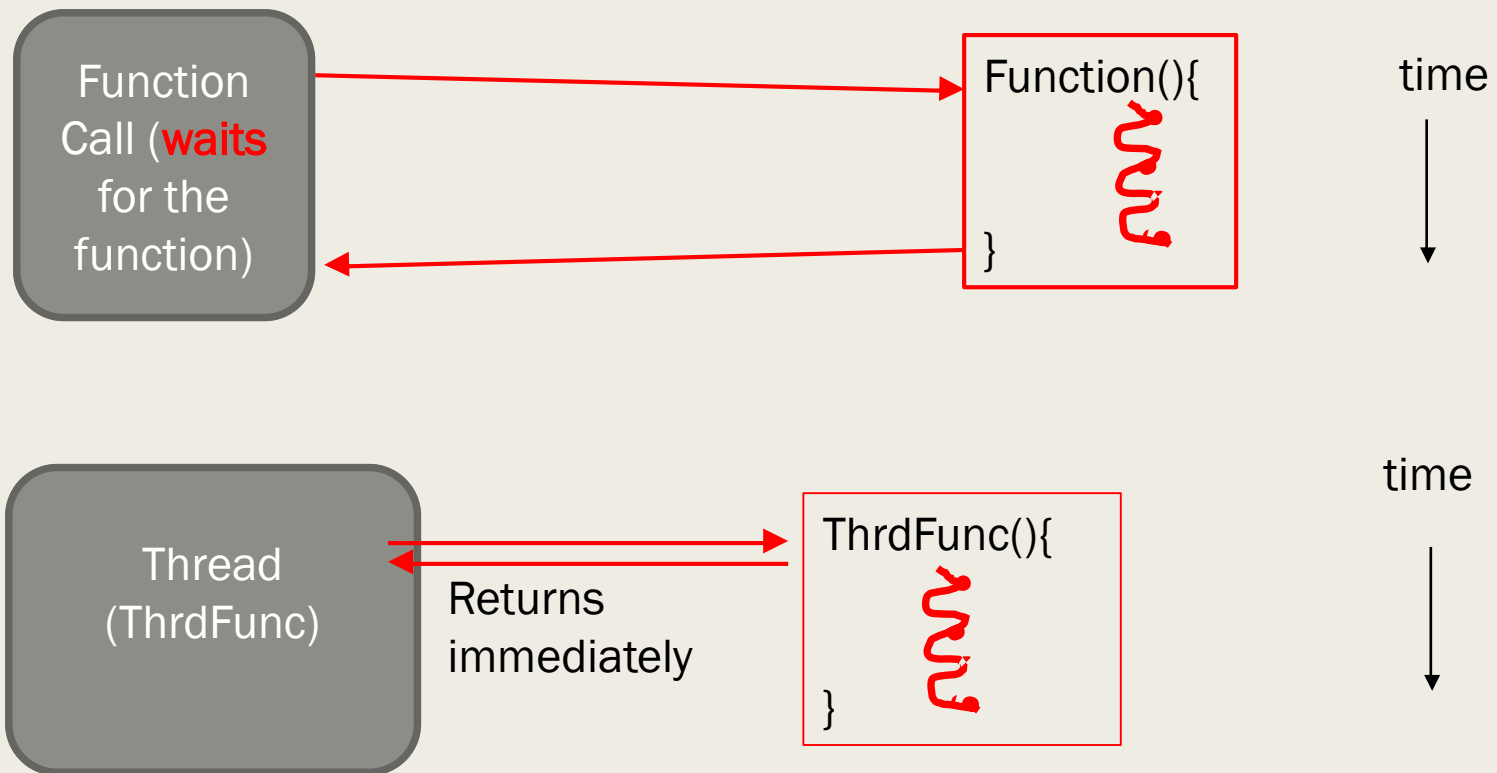
```
int main() {
    thread foothr (foo); //calls foo() in a thread
    thread barthr (bar,0);//calls bar(x) in a thread

    cout << "main, foo and bar would now execute
    concurrently..." << endl;;

    // synchronize threads:
    foothr.join(); // pauses until foo finishes
    barthr.join(); // pauses until bar finishes
    cout << "foo and bar completed.\n";
    return 0;
}
```

Thread Execution

- Creating a thread is very similar to calling a function directly, with a slight difference shown below:
 - *A regular function call is blocking, i.e., caller waits for callee*
 - *thread does **NOT***
- So, this is like `fork()`, creates the thread, calls the function inside, but does not wait



Race Condition Demonstration

- Start 2 (or more) threads that increment a shared variable times
 - *Use pointer to data so that threads share*
- Use a large number for x to ensure adequate overlap
- Notice that the output is different every time

```
void func (int * p, int x) {  
    // increment *p x times  
    for (int i=0; i<x; i++)  
        *p = *p + 1;  
}  
int main(int ac, char** av) {  
    int data = 0;  
    int times = atoi (av [1]);  
    // start 2 thread to increment the same variable  
    thread t1 (func, &data, times);  
    thread t2 (func, &data, times);  
  
    t1.join(); // pauses until first finishes  
    t2.join(); // pauses until second finishes  
    cout << "data = " << data << endl;  
    return 0;  
}
```

```
osboxes@osboxes:~/ $ ./a.out 10000000  
data = 15003189  
osboxes@osboxes:~/ $ ./a.out 10000000  
data = 13380972  
osboxes@osboxes:~/ $ ./a.out 10000000  
data = 12565682
```