



INTER-PROCESS COMMUNICATION

Tanzir Ahmed
CSCE 313 Spring 2020

Inter-Process Communication

■ IPC Methods

- *Pipes and FIFO*
- *Message Passing*
- *Shared Memory*
- *Semaphore Sets*
- *Signals*


■ References:

- *Beej's guide to Inter Process Communication for the code examples* (<https://beej.us/guide/bgipc/>)
- *Understanding Unix/Linux Programming, Bruce Molay, Chapters 10, 15*
- [*Advanced Linux Programming Ch 5*](#)
- *Some material also directly taken or adapted with changes from [Illinois course in System Programming](#) (Prof. Angrave), [UCSD](#) (Prof. Snoeren), and [USNA](#) (Prof. Brown)*

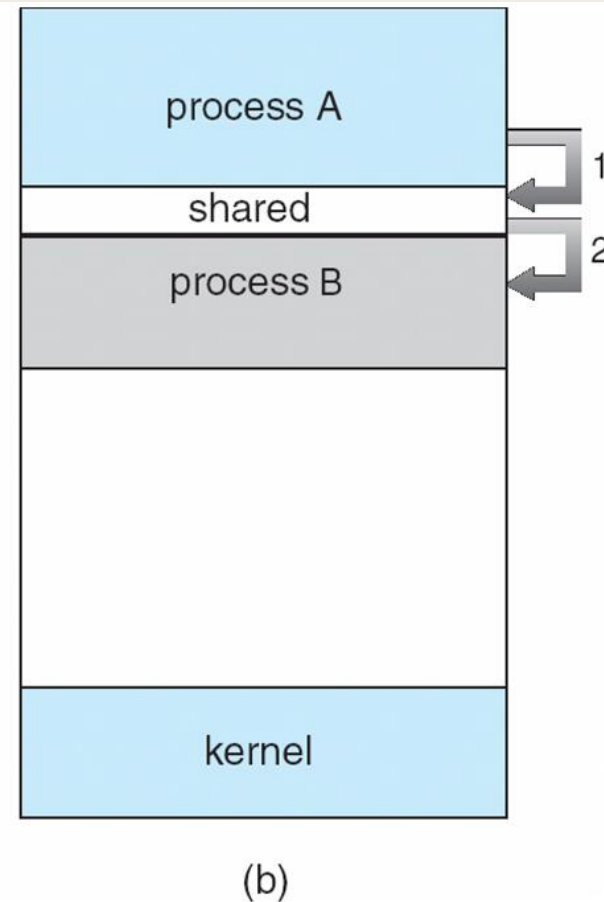
Inter-Process Communication

- OS provides generic mechanisms to communicate
 - *Un-named and Named Pipes*
 - Explicit communication channel
 - Explicit Synchronization: read() operation is blocking
 - *Message Passing: explicit communication channel provided through send()/receive() system calls*
 - Explicit Synchronization through send() and recv()
 - *Shared Memory: multiple processes can read/write same physical portion of memory; implicit channel*
 - Implicit channel
 - No OS mediation required once memory is mapped
 - No synchronization between communicating threads
 - *There is no read/send or write/recv available*
 - *How to know when data is available?*

IPC Fundamental Communication Models




Example: pipe, fifo, message, signal



Example: shared memory, memory mapped file

Communication Over a Pipe



Unix Pipes (aka Unnamed Pipes)

```
int pipe(int fildes[2])
```

- Returns a pair of file descriptors
 - *fildes[0] is connected to the read end of the pipe*
 - *fildes[1] is connected to the write end of the pipe*
- Create a message pipe
 - *Data is received in the order it was sent*
 - *OS enforces mutual exclusion: only one process at a time*
 - *Processes sharing the pipe must have same parent in common*


Pipe Creation

BEFORE pipe



Process has some usual files open

AFTER pipe



Kernel creates a pipe and sets file descriptors

- BEFORE
 - *Shows standard set of file descriptors*
- AFTER
 - Shows newly created pipe in the kernel and the two connections to that pipe in the process

IPC Pipe - Method

```
#include <stdio.h>
#include <unistd.h>


void main ()
{
    char buf [10];
    int fds [2];
    pipe (fds);
    printf ("sending msg: Hi\n");
    write (fds[1], "Hi", 3);
    read (fds[0], buf, 3);
    printf ("Received msg: %s\n", buf);
}
```

Connects the
two fds as pipe

```
compute-linux1 tanzir/code> ./a.out
sending msg: Hi
Received msg: Hi
```


Unnamed Pipe Between Two Processes

```
int main ()
{
    int fds [2];
    pipe (fds); // connect the pipe
    if (!fork()){ // on the child side
        sleep (3);
        char * msg = "a test message";
        printf ("CHILD: Sent %s\n", msg);
        write (fds [1], msg,
strlen(msg)+1);
    }else{
        char buf [100];
        read (fds [0], buf, 100);
        printf ("PRNT: Recvd %s\n", buf);
    }
    return 0;
}
```




Shell Piping: “ls -l | wc -l”

```
void main ()
{
    int fds [2];
    pipe (fds); // connect the pipe
    if (!fork()){ // on the child side
        close (fds[0]); // closing unnecessary pipe end
        dup2 (fds[1], 1); // overwriting stdout with pipeout
        execlp ("ls", "ls", "-l", NULL);
    }else{
        close (fds[1]); // closing unnecessary pipe end
        dup2 (fds[0], 0); // overwrite stdin with pipe in
        execlp ("wc", "wc", "-l", NULL);
    }
}
```

IPC- FIFO (named PIPE)

FIFO

- A pipe (also called unnamed pipe) disappears when no process has it open
- FIFOs (also called named pipes) are a mechanism that allow for IPC that's like using regular files, except that the kernel takes care of synchronizing reads and writes, and
- Data is never actually written to disk (instead it is stored in buffers in memory) so the overhead of disk I/O (which is huge!) is avoided.



FIFO vs PIPE

- A FIFO is like an unconnected garden hose lying on the lawn
 - *Anyone can put one end of the hose to his ear and another person can walk up to the hose and speak into the other end*
 - *Unrelated people may communicate through a hose*
 - *Hose exists even if nobody is using it*



FIFO

- **It's part of the file system**

- *It has a name and path just like a regular file.*
- *Programs can open it for reading and writing, just like a regular file.*
- *However, the file does not contain any data*

- **Works like a Bounded Buffer**

- *Bytes travel in First-In-First-Out fashion: hence the name FIFO.*
- *Reading process must wait for the writer*
- *Writer must wait for the read operation once the buffer is full*

FIFO - Problems

- We still need to agree on a name ahead of time – how to communicate that??

```
FIFORequestChannel rc ("control", ..) {  
    ...  
    mkfifo ("control", PERMS); // create  
}
```

- Not concurrency safe within a process
 - *Like a file used by multiple processes/threads*
 - *Multiple threads writing can cause race condition*

Using FIFO's

- How do I create a FIFO
 - *mkfifo (name)*
- How do I remove a FIFO
 - *rm fifoname or unlink(fifoname)*
- How do I listen at a FIFO for a connection
 - *open (fifoname, O_RDONLY)*
- How do I open a FIFO in write mode?
 - *open(fifoname, O_WRONLY)*
- How do two processes speak through a FIFO?
 - *The sending process uses write and the listening process uses read. When the writing process closes, the reader sees end of file*

FIFO DEMO

Writer

```
#define FIFO_NAME "test.txt"
int main(void)
{
    char s[300];
    int num, fd;
    mkfifo(FIFO_NAME, 0666); // create
    printf("Waiting for readers...\n");
    fd = open(FIFO_NAME, O_WRONLY); //open
    if (fd < 0)
        return 0;
    printf("Got a reader--type some
stuff\n");
    while (gets(s)) {
        if (!strcmp (s, "quit")) break;
        if ((num = write(fd, s, strlen(s)))
== -1)
            perror("write");
        else
            printf("SENDER: wrote %d bytes\n",
num);
    }
    //unlink (FIFO_NAME);
    return 0;
}
```

Reader

```
int main(void)
{
    char s[300];
    int num, fd;
    printf("waiting for writers...\n");
    fd = open(FIFO_NAME, O_RDONLY);
    printf("got a writer\n");
    do{
        if ((num = read(fd, s, 300)) == -1)
            perror("read");
        else {
            s[num] = '\0';
            printf("RECV: read %d bytes:
\"%s\"\n", num, s);
        }
    } while (num > 0);
    return 0;
}
```


IPC: Message Queue

Message Queue


- Mechanism for processes to communicate and to synchronize their actions
- IPC facility provides two operations:
 - ***send(message)***
 - ***receive(message)***
- If P and Q wish to communicate, they need to:
 - *establish a communication link between them*
 - *exchange messages via send/receive*

Message Passing

Direct



Indirect




Direct Message Passing

- Processes must name each other explicitly:
 - ***send*** (*P, message*) – send a message to process *P*
 - ***receive***(*Q, message*) – receive a message from process *Q*
- Properties of communication link
 - *Links are established automatically (or implicitly) while sending/receiving*
 - *A link is associated with exactly one pair of communicating processes*
 - *Between each pair, there exists exactly one link*
- Limitation: Must know the name or id of the other process



Indirect Message Passing

- Messages are directed to and received from **mailboxes** (also referred to as ports)
 - Mailbox can be owned by a process or by the OS
 - Each mailbox has a unique id
 - *Processes can communicate only if they share a mailbox*
- Properties of communication link
 - Link established only if processes share a **common mailbox**
 - A link may be associated with many processes
 - Each pair of processes may share several communication links



Indirect Message Passing


- Operations

- *create a new mailbox*
- *send and receive messages through mailbox*
- *destroy a mailbox*

- Primitives are defined as:

send(A, message) – send a message to mailbox A

receive(A, message) – receive a message from mailbox A



Synchronization

- Message passing may be either blocking or non-blocking
- **Blocking** is considered **synchronous**
 - ***Blocking send** has the sender block until the message is received*
 - ***Blocking receive** has the receiver block until a message is available*
- **Non-blocking** is considered **asynchronous**
 - ***Non-blocking send** has the sender send the message and continue*
 - ***Non-blocking receive** has the receiver receive a valid message or null*
 - Does not wait for it
 - A receive may “fail” several times before it succeeds at the end

Operations on Message Queues

```
mqd_t mq_open(const char *name, int oflag,  
              mode_t mode, struct mq_attr *attr);
```

```
int mq_send(mqd_t mqdes, const char *msg_ptr,  
            size_t msg_len, unsigned int msg_prio);
```

```
ssize_t mq_receive(mqd_t mqdes, char *msg_ptr,  
                   size_t msg_len, unsigned int *msg_prio)
```

```
int mq_close(mqd_t mqdes)
```

Message Queue – Example

```
send() {  
    mqd_t mq = mq_open("/testqueue", O_RDWR|O_CREAT, 0664, 0);  
    if (mq_send(mq, av[1], strlen (av[1]) + 1, 0)<0){  
        perror ("MQ Send failure");  
        exit (0);  
    }  
    printf ("MQ Put: %s\n", av [1]);  
    return 0;  
}
```

```
recieve() {  
    mqd_t mq = mq_open("/testqueue",O_RDWR|O_CREAT, 0664,0);  
    struct mq_attr attr;  
    mq_getattr (mq, &attr); // get attribute  
    char *buf = (char*)malloc (attr.mq_msgsize);  
    mq_receive(mq, buf, attr.mq_msgsize, NULL);  
    printf ("MQ Receive Got: %s\n", buf);  
    //clean-up  
    mq_close(mq);  
    //mq_unlink("/testqueue"); // remove from Kernel  
    return 0;  
}
```

IPC: Shared Memory


Shared Memory

- How does data travel through a FIFO?
 - *‘write’ copies data from process memory to kernel buffer and then ‘read’ copies data from a kernel buffer to process memory*
- If both processes are on the same machine, then they may not need to copy data in and out of the kernel
 - *They may exchange or share data by using a shared memory segment*
 - *Shared memory is to processes what global variables are to threads*

Shared Memory

- Processes share the same segment of memory directly
 - *Memory is mapped into the address space of each sharing process*
 - *Memory is persistent beyond the lifetime of the creating or modifying processes (until deleted)*
- But, now the processes are on their own for synchronization
 - ***Mutual exclusion must be provided by processes using the shared memory***
 - ***Semaphores*** – *to ensure Producer-Consumer relation also has to be now done by the respective processes*

Shared Memory



- Processes request the segment
- OS maintains the segment
- Processes can attach/detach the segment

Facts about Shared Memory Segments

- A shared memory segment has a name, the name must start with “/”
- A shared memory segment has an owner and permission bits
- Processes may “map” or “unmap” a segment, obtaining/removing a pointer to the segment
- reads and writes to the memory segment are done via regular pointer operations

Shared Memory – POSIX functions

- **shm_open**: creates a shared memory segment
- **ftruncate**: sets the size of a shared memory segment
- **mmap**: maps the shared memory object to the process's address space
- **munmap**: unmaps from process's address space
- **shm_unlink**: removes the shared memory segment from the kernel
- **close**: closes the file descriptor associated with the shared memory segment

Shared Memory Example

```
char* my_shm_connect(char* name, int len){
    int fd = shm_open(name, O_RDWR|O_CREAT, 0644 );
    ftruncate(fd, len); //set the length to 1024, the default
    is 0, so this is a necessary step
    char *ptr = (char *) mmap(NULL, len, PROT_READ|PROT_WRITE,
MAP_SHARED, fd, 0); // map
    if (fd < 0){
        perror ("Cannot create shared memory\n");
        exit (0);
    }
    return ptr;
}

void send(char* message){
    char *name = "/testing";
    int len = 1024;
    char* ptr = my_shm_connect (name, len);

    strcpy(ptr, message); // putting data by just copying
    printf ("Put message: %s\n", message);
    close(fd); // close desc, does not remove the segment
    munmap (ptr, len); // this is a bit redundant,
}
```

Shared Memory Example- contd

```
void receive(){
    char    *name    = "/testing";
    int len = 1024;
    char* ptr = my_shm_connect (name, len)

    printf ("Got message: %s\n", ptr);
    shm_unlink (name); //this removes the segment
from Kernel, this is a necessary clean up
    exit(0);
}
```

Shared Memory - Summary

- It's great because it does not allocate extra memory
 - *3-times less memory (In others you have 1 copy in kernel, 2 others in the individual processes)*
- But a huge problem looms
 - *Who does synchronization?*
 - *There is no guarantee of order between sending and receiving processes*
 - *The receiver process can run first finding nothing or stale data in the buffer*

Kernel Semaphores

- We have learned how to synchronize multiple threads using Semaphores and Locks
- But how do we synchronize multiple processes?
 - *We will again need semaphores, but this time Kernel Semaphores*
 - *They are visible to separate processes who do not share address space*

Kernel Semaphore Operations

- **sem_open(name, ...)** to create or connect to a semaphore
 - *The name argument must start with a "/"*
- **sem_close ()** closes a semaphore
 - *It does not destroy it from Kernel*
- **sem_unlink()** removes from Kernel
 - *Must be put in the destructor for PA5*
- **sem_wait()** is equivalent to Semaphore::P() operation
- **sem_post()** is equivalent to Semaphore::V() operation
- Find out more from **sem_overview(7)** in man pages